

Lab 01 — Meet ModSecurity & the OWASP Core Rule Set

Level: Beginner **Goal:** Understand what a WAF is, stand up the lab stack, and watch your very first request get either *allowed* or *blocked*. **Time:** ~20 minutes

1. Theory first — what problem are we solving?

As a web developer you already fix vulnerabilities *in code* — you escape output, you use parameterised SQL queries, you validate input. That is the right place to fix things. So why add a WAF?

Because in the real world:

- You don't control every line of code (third-party libraries, legacy endpoints, a plugin someone installed).
- A brand-new vulnerability (a "0-day" / fresh CVE) can be published on a Monday and exploited worldwide by Tuesday — long before you can patch, test and redeploy.
- Some bugs live in the framework or the runtime itself.

A **WAF (Web Application Firewall)** sits *in front of* your app and inspects every HTTP request **before your code ever runs**. If a request looks like an attack, the WAF can reject it. Think of it as a bouncer at the door who reads each request and turns away the obviously malicious ones.



The three pieces

Piece	What it is	Analogy
ModSecurity	An open-source <i>engine</i> that can inspect HTTP requests/responses and act on them. On its own it has no rules — it just knows <i>how</i> to match rules.	The bouncer's brain
OWASP Core Rule Set (CRS)	A community-maintained <i>set of rules</i> describing what attacks look like (SQLi, XSS, path traversal, etc.). This is what makes ModSecurity actually useful.	The bouncer's rulebook
nginx (or Apache)	The web server ModSecurity plugs into. The <code>owasp/modsecurity-crs:nginx</code> image bundles all three, pre-wired.	The doorway itself

Why this matters: People say "I installed ModSecurity" and think they are protected. ModSecurity **without** a rule set blocks nothing. The protection comes from the **CRS**. The Docker image you already have bundles both, which is exactly why we use it.

2. The lab stack

Everything in these labs uses two containers:

- **app** — a tiny, deliberately naive Flask web app (the thing we protect). It has endpoints that *would* be vulnerable to XSS, SQLi and path traversal if nothing stood in front of it.
- **waf** — the `owasp/modsecurity-crs:nginx` container. It listens on `:8080`, runs each request through the CRS, and only forwards clean requests to the Flask app.

You attack `http://localhost:8080` (through the WAF). If you ever want to see the *unprotected* app for comparison, it is on `http://localhost:5000`.

3. Practical steps

Step 3.1 — Get the image (you already have it)

```
docker pull owasp/modsecurity-crs:nginx
docker images | grep modsecurity
```

Step 3.2 — Start the whole stack

From the repository root:

```
docker compose up -d --build
```

This builds the Flask app image and starts both containers. Check they are up:

```
docker compose ps
```

You should see `modseclabs-app` and `modseclabs-waf` running.

Step 3.3 — Send your first (clean) request

```
curl -i http://localhost:8080/
```

You get back `HTTP/1.1 200 OK` and the app's HTML. The request was clean, so the WAF passed it through. In a browser it looks like this:

Clean request through the WAF -> reaches Flask app [HTTP 200]

 ModSec Labs – behind OWASP CRS

Target Web Application

This tiny app is the origin server the labs attack. If you can read this through `http://localhost:8080`, your request passed the WAF.

1. Reflected search (XSS playground)

Search

2. User lookup (SQL-injection playground)

Look up

3. File viewer (path-traversal playground)

Open

Endpoints: `/search?q=` · `/login?user=` · `/file?name=` · `/api/log` (POST) · `/healthz`

Step 3.4 — Send your first attack

Now ask for the same app but smuggle a classic cross-site-scripting payload in the `q` parameter:

```
curl -i "http://localhost:8080/search?q=<script>alert(1)</script>"
```

This time the response is **HTTP/1.1 403 Forbidden** — the WAF recognised the `<script>` payload and refused to forward it. Your Flask code never ran.

```
GET /search?q=<script>alert(1)</script> -> BLOCKED by CRS 941 [HTTP 403]
```

403 Forbidden

nginx

Where the theory shows up: The **403** did **not** come from your Flask app — Flask never saw the request. It came from ModSecurity applying a CRS rule at the front door. That is the whole idea of a WAF in one screenshot.

Step 3.5 — Prove the app itself is defenceless

To really appreciate the WAF, hit the app **directly**, bypassing the WAF:

```
curl -i "http://localhost:5000/search?q=<script>alert(1)</script>"
```

Now you get **200 OK** and your **<script>** tag is reflected right back into the HTML — a working reflected-XSS. Same app, same payload; the only difference is whether the WAF was in the path.

4. What just happened, rule by rule

When the attack was blocked, ModSecurity wrote an **audit log** entry. View it:

```
docker logs modseclabs-waf | tail -n 40
```

Buried in there you'll find the exact rule that fired:

```
id "941100" msg "XSS Attack Detected via libinjection"
id "949110" msg "Inbound Anomaly Score Exceeded (Total Score: 20)"
```

Two important ideas are already visible here, and we'll go deep on both in later labs:

- **941100** — a rule from the XSS family (**941xxx**) spotted the attack.
- **949110** — CRS didn't block on the first hit; it *added up* a score and blocked once the total crossed a threshold. This is **anomaly scoring** (Lab 02 and Lab 05).

5. Clean up (optional)

```
docker compose down
```

6. Recap

- A **WAF** inspects HTTP traffic *before your code runs*.
- **ModSecurity = engine**, **CRS = rules**. You need both; the Docker image bundles both.
- A clean request → **200** ; a recognised attack → **403** , and your app never sees it.
- Every block is explained in the **audit log**, tied to a numbered rule.

Next: In Lab 02 we stop *blocking* and switch the WAF into **DetectionOnly** mode, learn to read the audit log properly, and understand the anomaly score that decided the block.