

Lab 02 — Detection vs Prevention & Reading the Audit Log

Level: Beginner **Goal:** Understand the three engine modes, learn to read a ModSecurity audit log, and understand the **anomaly score** that decides a block. **Time:** ~30 minutes

1. Theory — the rule engine has three modes

ModSecurity's behaviour is controlled by one master switch, the `SecRuleEngine` directive. It has three settings:

Mode	What it does	When you use it
<code>On</code>	Rules run and requests can be blocked.	Production, once you trust your tuning.
<code>DetectionOnly</code>	Rules run and everything is logged , but nothing is ever blocked.	Always start here. Watch what <i>would</i> be blocked before you turn on blocking.
<code>Off</code>	The engine is disabled. No inspection at all.	Debugging / temporarily bypassing.

Why this matters — the #1 rule of deploying a WAF: Never switch a WAF straight to `On` in front of a real app. You *will* have false positives (legitimate requests that look suspicious). Run in `DetectionOnly` first, read the logs, fix the false positives (Lab 07), and only then enable blocking. Flipping straight to `On` is how teams take down their own site on day one.

In the lab stack the mode is set by the `MODSEC_RULE_ENGINE` environment variable, which the container maps to `SecRuleEngine`.

2. Practical — see the difference with your own eyes

We'll run the **same attack** against two WAFs that differ only in mode.

Step 2.1 — Start a DetectionOnly WAF alongside the blocking one

The main stack (`:8080`) runs in `On` mode. Start a second WAF in `DetectionOnly` mode on `:8081` :

```
docker run -d --name modseclabs-waf-detect \  
  --add-host=host.docker.internal:host-gateway \  
  -e BACKEND="http://host.docker.internal:5000" -e PORT=8080 \  
  -e MODSEC_RULE_ENGINE=DetectionOnly -e PARANOIA=1 -e ANOMALY_INBOUND=5 \  
  -e ANOMALY_SCORE_THRESHOLD=5
```

```
-e MODSEC_AUDIT_LOG=/dev/stdout -e MODSEC_AUDIT_ENGINE=RelevantOnly \  
-p 8081:8080 owasp/modsecurity-crs:nginx
```

In `docker compose`, you'd instead just change `MODSEC_RULE_ENGINE` on the `waf` service and `docker compose up -d`. We use a second container here so you can compare both live.

Step 2.2 — Fire the same XSS payload at both

```
# Blocking mode  
curl -s -o /dev/null -w "Prevention (:8080) -> HTTP %{http_code}\n" \  
"http://localhost:8080/search?q=<script>alert(1)</script>"  
  
# Detection-only mode  
curl -s -o /dev/null -w "DetectionOnly (:8081) -> HTTP %{http_code}\n" \  
"http://localhost:8081/search?q=<script>alert(1)</script>"
```

You'll see:

```
Prevention (:8080) -> HTTP 403  
DetectionOnly (:8081) -> HTTP 200
```

Same rule fires in both. The difference is only the **consequence**. In `DetectionOnly` the malicious request sails through to the app:

DetectionOnly mode: attack PASSES (HTTP 200) but is logged [HTTP 200]

Results for:

No results found.

Where the theory shows up: Both engines *detected* the attack — check the logs of either and rule **941100** fired. Detection ≠ prevention. The engine mode is the switch between "know about it" and "stop it".

3. Anatomy of an audit log entry

This is the single most useful skill in operating a WAF. When something is blocked (or would be), ModSecurity writes a JSON audit record. View it:

```
docker logs modseclabs-waf | tail -n 40
```

The record is big, but you only need a handful of fields. Here are the ones that matter, from a real block in this lab:

```
"ruleId": "941100",  
"msg": "XSS Attack Detected via libinjection",  
"data": "Matched Data: XSS data found within ARGS:q: <script>alert(1)</script>",
```

```
"severity": "2",
"tags": ["attack-xss", "paranoia-level/1", "OWASP_CRS/ATTACK-XSS"]
```

Field	Meaning	Why you care
<code>ruleId</code>	The numeric ID of the rule that fired (<code>941100</code>).	This is the handle you use to tune, disable, or research a rule.
<code>msg</code>	Human description of what was detected.	Tells you the attack class.
<code>data</code>	Exactly what matched, and in which variable (<code>ARGS:q</code>).	Lets you confirm true vs false positive at a glance.
<code>severity</code>	How serious CRS thinks it is (lower number = worse).	Feeds the anomaly score.
<code>tags</code>	Categories incl. the paranoia level that caught it.	Used for bulk exclusions and reporting.

Rule ID families (memorise these ranges)

CRS groups rules by attack type using the leading digits of the ID:

Range	Attack family
<code>913xxx</code>	Scanner / bot detection
<code>920xxx – 921xxx</code>	Protocol enforcement (malformed HTTP)
<code>930xxx</code>	Local File Inclusion / path traversal
<code>931xxx</code>	Remote File Inclusion
<code>932xxx</code>	Remote Code Execution (shell injection)
<code>941xxx</code>	Cross-Site Scripting (XSS)
<code>942xxx</code>	SQL Injection
<code>943xxx</code>	Session fixation
<code>944xxx</code>	Java attacks (incl. Log4Shell — see Lab 08)
<code>949xxx</code>	Blocking evaluation (the anomaly-score decision)

4. Theory — anomaly scoring (the heart of CRS)

Look again at the Lab 01 block. It contained **two** rule hits:

```
941100 XSS Attack Detected (severity CRITICAL -> +5)
949110 Inbound Anomaly Score Exceeded (Total Score: 20)
```

CRS 4.x does **not** block the instant one rule matches. Instead:

1. Every matching rule **adds points** to a running *anomaly score*. A **CRITICAL** rule adds 5, **ERROR** adds 4, **WARNING** 3, **NOTICE** 2.
2. At the very end (rule **949110**), CRS compares the total against the **inbound anomaly threshold** (**ANOMALY_INBOUND**, default 5).
3. If **score >= threshold**, the request is blocked. Otherwise it passes.

```
payload matches 941100 (+5), 941110 (+5), 941160 (+5), 941390 (+5)
      |
      v
total inbound score = 20
      |
949110: is 20 >= threshold(5) ? YES -> 403 Forbidden
```

Why this design matters: A single fuzzy rule firing on its own might be a false positive. But a request that trips *four* XSS rules is almost certainly an attack. Scoring lets CRS be **strict about real attacks and forgiving about borderline traffic** by tuning **one number** (the threshold) instead of editing hundreds of rules. You'll tune this threshold and the "paranoia level" that changes how many rules run in Lab 05.

The threshold is a dial:

- **Lower threshold (e.g. 3)** → blocks more aggressively → more false positives.
- **Higher threshold (e.g. 10)** → more permissive → risks missing real attacks.

5. Exercises

1. Fire the SQLi payload `admin' OR '1'='1` at `:8081` and confirm it returns `200` but still logs rule `942100`.
2. In the audit log, find the `data` field and confirm it shows *your* exact payload. This is how you distinguish a real attack from a false positive.
3. Change the main stack to **DetectionOnly** (edit `MODSEC_RULE_ENGINE` in `docker-compose.yml`, then `docker compose up -d`) and confirm the XSS now returns `200`.

Clean up the extra container

```
docker rm -f modseclabs-waf-detect
```

6. Recap

- `SecRuleEngine` : **On** (block), **DetectionOnly** (log only), **Off** (disabled). **Always start in DetectionOnly.**
- The audit log's `ruleId`, `msg` and `data` fields tell you *what* fired and *why* — learn to read them.

- Rule ID ranges map to attack families (941 = XSS, 942 = SQLi, ...).
- CRS blocks on an **accumulated anomaly score** crossing a threshold, not on a single rule hit.

Next: Lab 03 dives into a full attack family — SQL injection — and shows exactly which payloads trip which rules. 942xxx